CASE STUDY

AI pipelines that run real work, with a person approving every output

Six production pipelines for project management and operations. Each has its own job, its own schedule and its own approval step. Nothing sends, publishes or deletes without a person signing off.

Built for my own work and for clients | A person approves every output | Measured counts, estimated hours

At a glance

Over 400
meetings prepared and followed up

Over 300
candidate applications assessed

Over 200
weekly project reports drafted

Over 2,500
job postings scored

The pipelines

Transcript and action-point follow-up

Meeting transcripts become a searchable record: summary, decisions, owners and priorities. Every action item is logged, turned into a task and kept on the reports until it is closed.

Over 400 meetings, over 200 of them captured in the database. Over 1,300 action points, over 700 closed. Run for five clients.

By hand: 45 to 60 minutes of prep and 45 to 60 of notes per meeting, and the notes caught 5 to 10% of what was said. Now: 10 to 15 minutes of prep and about 20 of review. Estimated: over 400 hours saved.

Talent pipeline

Reads every application, scores it against set criteria, reviews the candidate's code, files the CV and drafts a personal reply with advice on how to improve. A resubmission is scored again.

Over 300 applications. Over 100 scored and filed. Over 50 code reviews. Over 30 interviewed. 8 placed, and the first 5 have moved on to full-time roles. All figures are annual. About 90% automated. Over half came back at least once.

Personal details are stripped before anything reaches the model, and every applicant is told by email that AI is used. By hand: 60 to 90 minutes per candidate plus 15 to 30 for the task card, and code review never happened. Now: about 30 minutes of my review. Estimated: over 200 hours saved.

Weekly project reports

One report per project, drafted from the record. I edit it and send it.

Over 200 reports.

By hand: 60 to 120 minutes of looking up, documenting and writing, and the report got skipped when the workload was high. Now it gets edited and sent. Estimated: over 200 hours saved.

Digests

Priorities from every source combined into one view that surfaces what matters most, without opening email or being pulled off task.

Teams chatbot

Ask in plain English what was decided, who owns what, or where an applicant stands. It answers from the database and sends updates on open items. Used daily.

Job sourcing

Scores every new job posting against a written skill profile, three times a day, and ranks them. I decide what to apply for.

Over 2,500 postings scored.

Nobody reads that many by hand.

The talent pipeline, step by step

Built as a development pipeline, not a filter. Nobody is ghosted: every applicant gets a reply, a score and specific advice, and can come back. Each round shows who is most motivated and who is most skilled.

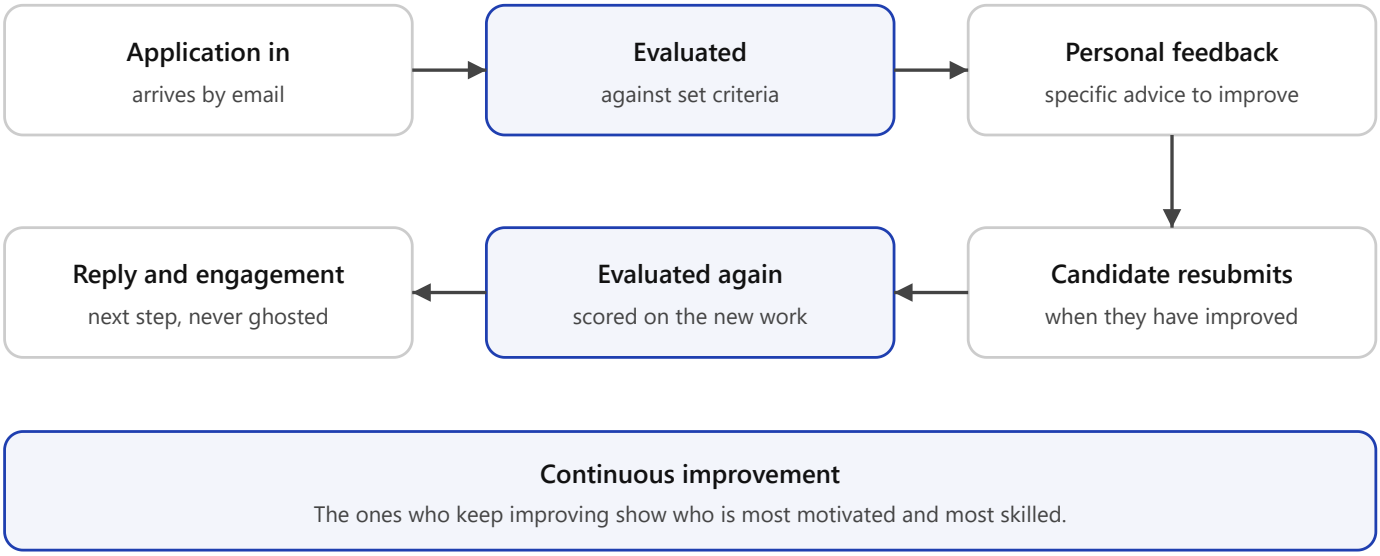
Over 300
applications received in a year

Over 30
candidates interviewed in a year

8
placed in a year, writing production code

5
of the first interns have moved on to full-time roles

About 90% of the process is automated, with a person in the loop for the last 10%. One person runs it, in production for over a year. The earlier manual version, at another organization, took three people four years to place 6 to 8 students a year. The 90% is my estimate. Two new students join at every break, about 8 a year. The figures above are annual.



What it looks like

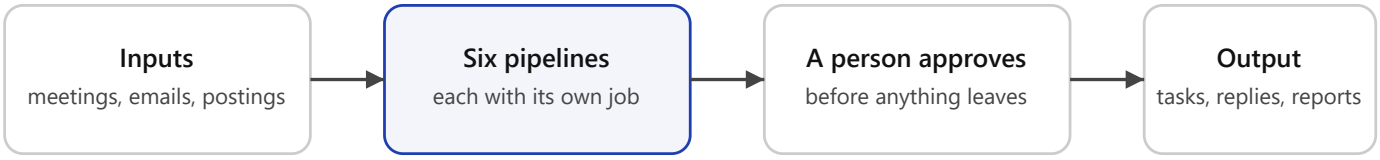


The approval screen for meeting action points. Every action is listed with owner, project and priority, and nothing is published to Planner until a person approves it. Names and details are blurred.



The top of the daily digest: urgent work, blocked work, drafts waiting and strong candidates at a glance. The rest of the digest is not shown.

How they work



- A person approves before anything is sent, published or changed.
- Fixed rules own the numbers. AI handles only what rules cannot.
- Every automated action is logged.
- They read and write in Microsoft 365 (Teams, Planner, SharePoint) through Microsoft Graph.
- Each task goes to the right model, which cut AI cost by about 70%.

How to read the numbers

Counts come from the pipelines' own records. Hours are estimates: the minutes the work took by hand before, against the minutes of review now, taking the low end of each range. Where I would never have done the work by hand, such as capturing a whole meeting or reading every posting, I count it as capacity, not hours.

Where it is now

All six run today, for my own work and for clients. The transcript and action-point follow-up has run for five separate clients.

Terms: AI automation, AI agents, workflow automation, meeting intelligence, action-item tracking, candidate screening, project reporting, chatbot, Microsoft Graph, human in the loop, GDPR.

Mark Tenzer, technical project manager and integration specialist. LinkedIn

Case study v0.2, 1 November 2026. Names and branding removed. No tracking and no outside scripts.